

Lecture 10: Transactions & Concurrency Control

Dr. Kyong-Ha Lee
(kyongha@kisti.re.kr)









Contents

Brief overview of this lecture

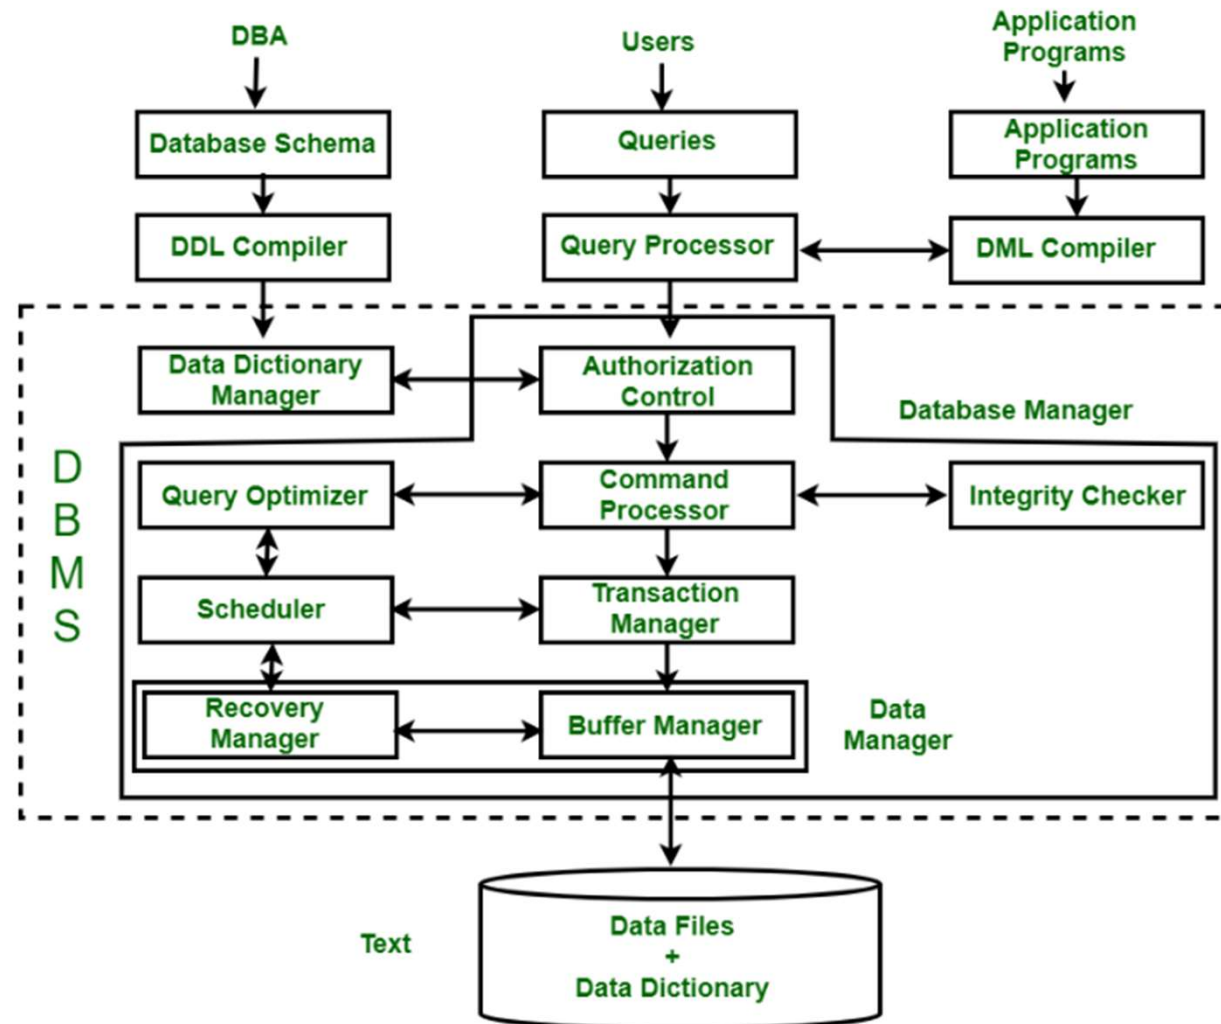
- Basic theories and principles about Index mechanisms used in database systems
- Not much discussion on implementation or tools, but will be happy to discuss them if there are any questions

 Contents 1	Transaction Concepts
 Contents 2	Concurrent Execution
 Contents 3	Serializability
 Contents 4	Implementation of Isolation

*Disclaimer: these slides are based on the slides created by the authors of Database System Concepts 7th ed. and modified by K.H. Lee.



DBMS Architecture Overview





Transaction

- A **unit** of execution that accesses and possibly update various data items
 - From the DBMS's point of view, a transaction is a series of the following actions
 - **READs:** DB object is read from disk into buffer page.
 - **WRITEs:** DB object is written from buffer to disk
 - **ABORT:** Last action of a Xact that fails
 - **COMMIT:** Last action of a Xact that succeeds
 - e.g., transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- Two main issues to deal with :
- Failures of various kinds, such as HW failures and system crashes
 - Concurrent execution of multiple transactions




Properties of Transactions (ACID)

- **Atomicity** : all or nothing
 - If system crashes, partially done transactions must be undone or "rolled back"
- **Consistency**: if each transaction is consistent and the DB starts consistent, it must be ended consistently
- **Isolation**: transactions are protected from the effects of all other concurrent transactions
 - Concurrency for performance must not affect each transaction.
- **Durability**: once the transaction has completed, updates by the transaction must persist even if SW or HW failures come



Transaction State

- **Active** – Initial state; TX stays in this state while it is executing
- **Partially committed** – after final statement has been executed
- **Failed** – after that normal execution can no longer proceed
- **Aborted** – after TX was rolled back and DB restored to its state prior to the start of the TX
 - Restart the TX
 - Kill the transaction
- **Committed** – after successful completion





Isolation requirement

- Isolation requirement
 - Isolation can be ensured trivially by running transactions **serially**. That is, one after the other.

T1

1. **read**(A)
2. $A := A - 50$
3. **write**(A)

4. **read**(B)
5. $B := B + 50$
6. **write**(B)

T2

read(A), read(B), print(A+B)

- However, executing multiple transactions concurrently has significant benefits



Concurrent Execution

- Users submit transactions and can think of each transaction as executing by itself (Give users such illusion)
 - Concurrency is achieved by the DBMS that **interleaves actions (reads/writes of DB objects) of various transactions**
 - Advantages
 - **Increased processor and disk utilization**, leading to better transaction *throughput*
 - E.g., one transaction can be using the CPU while another is reading from or writing to the disk
 - **Reduced average response time** for transactions: short transactions need not wait behind long ones.



Concurrent Execution

- Each transaction must leave the database in a consistent state if the DB is consistent when the transaction begins
 - DBMS will enforce some Integrity Constraints(IC), depending on the ICs declared in CREATE TABLE statements
 - Beyond this, the DBMS does not really understand the semantics of the data (e.g., it does not understand how the interest on a bank account is computed)
- Issues: effect of *interleaving* transactions and *system crashes*
- **Concurrency control schemes** – mechanisms to achieve isolation



Schedule

- **Schedule**

- a sequences of instructions that specify the **chronological order** in which instructions of concurrent transactions are executed
- A schedule for a set of transactions **must consist of all instructions** of those transactions
- Must **preserve the order** in which the instructions appear in each transaction.

- **Examble**

- Let T_1 transfer \$50 from A to B, and T_2 transfer 10% of the balance from A to B.
- A **serial** schedule in which T_1 is followed by T_2 :

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 1



- Let T_1 and T_2 be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1

T_1	T_2
read (A) $A := A - 50$	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	
	$B := B + temp$ write (B) commit

Schedule 2

T_1	T_2
read (A) $A := A - 50$ write (A)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	
	read (B) $B := B + temp$ write (B) commit

Schedule 3



Serializability

- **Basic Assumption** – Each transaction preserves database consistency.
 - serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule
 - Different forms of schedule equivalence give rise to the notions of:
 1. **Conflict serializability**
 2. **View serializability**

Note: we ignore operations other than read and write instructions



Conflicting instruction

- Instructions I_i and I_j of transactions T_i and T_j respectively, **conflict** if and only if there exists some item Q accessed by both I_i and I_j , and at least one of these instructions wrote Q .
 1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict.
 2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict
 4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict
- Conflict between I_i and I_j forces a (logical) temporal order between them.
- If I_i and I_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.



- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are **conflict equivalent**.
- We say that a schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule
- Schedule 3 is conflict serializable and schedule 4 is not. Why?
 - Unable to swap instructions in schedule 4 to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

T_1	T_2
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

Schedule 3

T_3	T_4
read (Q) write (Q)	write (Q)

Schedule 4



View Serializability

- Schedule S and S' are **view equivalent** if the 3 conditions are met, for each data item Q
 1. If in schedule S , transaction T_i reads the initial value of Q , then in schedule S' also transaction T_i must read the initial value of Q
 2. If in schedule S transaction T_i executes **read**(Q), and that value was produced by transaction T_j (if any), then in schedule S' also transaction T_i must read the value of Q that was produced by the same **write**(Q) operation of transaction T_j .
 3. The transaction (if any) that performs the final **write**(Q) operation in schedule S must also perform the final **write**(Q) operation in schedule S' .



View Serializability(Cont.)

- A schedule S is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Below is a schedule which is view-serializable but *not* conflict serializable.




T_{27}	T_{28}	T_{29}
read (Q)	write (Q)	write (Q)
write (Q)		

- What serial schedule is above equivalent to?
- Every view serializable schedule that is not conflict serializable has **blind writes**.



Testing for serializability

- **Precedence graph** — a direct graph where the vertices are the transactions (names)
 - Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
 - We draw an arc from T_i to T_j if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier.
 - We may label the arc by the item that was accessed.
- A schedule is conflict serializable if and only if its precedence graph is acyclic
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.





Recoverable Schedule

- We need to address the effect of transaction failures on concurrently running transactions
- **Recoverable schedule**
 - if a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i appears before the commit operation of T_j .
- The following schedule (Schedule 11) is not recoverable

T_8	T_9
read (A) write (A)	
	read (A) commit
read (B)	

- DB system must ensure recoverable schedules
 - If T_8 should abort, T_9 would have read (and possibly shown to the user) an inconsistent database state



Cascading Rollbacks

- **Cascading rollback**

- a single TX failure leads to a series of transaction rollbacks.
- Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

If T_{10} fails, T_{11} and T_{12} must also be rolled back.

- Can lead to the undoing of a significant amount of work

- **Cascadeless schedules**

- For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
 - Every Cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless



Concurrency Control

- A database must provide a mechanism that will ensure that all possible schedules are
 - either conflict or view serializable, and
 - are recoverable and preferably cascadeless
- Serial Schedule provides a poor degree of concurrency
- Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur.
- Some schemes allow only conflict-serializable schedules to be generated, while others allow view-serializable schedules that are not conflict-serializable.



Weak Levels of Consistency

- Some applications are willing to live with **weak levels of consistency**, allowing schedules that are not serializable
 - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - E.g., database statistics computed for query optimization can be approximate (why?)
 - Such transactions need not be serializable with respect to other transactions
- Tradeoff accuracy for performance



Levels of Consistency in SQL-92

- **Serializable** — default
- **Repeatable read** — only committed records to be read.
 - Repeated reads of same record must return same value.
 - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- **Read committed** — only committed records can be read.
 - Successive reads of record may return different (but committed) values.
- **Read uncommitted** — even uncommitted records may be read.



Levels of Consistency

- Lower degrees of consistency useful for gathering approximate information about the database
- Warning: some database systems do not ensure serializable schedules by default
- E.g., Oracle (and PostgreSQL prior to version 9) by default support a level of consistency called snapshot isolation (not part of the SQL standard)
- MariaDB offer all four consistency levels
 - Default isolation level for InnoDB is REPEATABLE READ



- 24



Implementation of Isolation Levels

- Locking
 - Lock on whole database vs lock on items
 - How long to hold lock?
 - Shared vs exclusive locks
- Timestamps
 - Transaction timestamp assigned e.g. when a transaction begins
 - Data items store two timestamps
 - Read timestamp
 - Write timestamp
 - Timestamps are used to detect out of order accesses
- Multiple versions of each data item
 - Allow transactions to read from a “snapshot” of the database



Question?

—source: <https://www.fox.com/the-simpsons>

